

CS 235: Artificial Intelligence

Week 3

Heuristic (Informed) Search

Dr. Moumita Roy
CSE Dept., IITG

Reference: <http://ai.stanford.edu/~latombe/cs121/2011/schedule.htm>

Recall that the ordering
of FRINGE defines the
search strategy

Search Algorithm #2

SEARCH#2

1. INSERT(initial-node, FRINGE)
2. Repeat:
 - a. If empty(FRINGE) then return failure
 - b. $N \leftarrow \text{REMOVE}(\text{FRINGE})$
 - c. $s \leftarrow \text{STATE}(N)$
 - d. If GOAL?(s) then return path or goal state
 - e. For every state s' in SUCCESSORS(s)
 - i. Create a node N' as a successor of N
 - ii. INSERT(N' , FRINGE)

Best-First Search

- It exploits **state description** to estimate how “good” each search node is
- An **evaluation function** f maps each node N of the search tree to a real number
 $f(N) \geq 0$
[Traditionally, $f(N)$ is an estimated cost; so, the smaller $f(N)$, the more promising N]
- **Best-first search** sorts the FRINGE in increasing f
[Arbitrary order is assumed among nodes with equal f]
- The strategy is identical to that for uniform cost search; except the use of f instead of g to order the priority queue.

How to construct f ?

- Typically, $f(N)$ estimates:

- either the **cost of a solution path through N**

Then $f(N) = g(N) + h(N)$, where

- $g(N)$ is the cost of the path from the initial node to N
- $h(N)$ is an estimate of the cost of a path from N to a goal node

- or the **cost of a path from N to a goal node**

Then $f(N) = h(N)$

A* search algorithm

Greedy best first search

Heuristic function

- But there are no limitations on f . Any function of your choice is acceptable.

But will it help the search algorithm?

Heuristic Function

- The **heuristic function** $h(N) \geq 0$ estimates the cost to go from $STATE(N)$ to a goal state

Its value is **independent of the current search tree**; it depends only on $STATE(N)$ and the goal test $GOAL?$

- Example:

5		8
4	2	1
7	3	6

$STATE(N)$

1	2	3
4	5	6
7	8	

Goal state

$h_1(N)$ = number of misplaced numbered tiles = 6

Other Examples

5		8
4	2	1
7	3	6

STATE(N)


1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$ = number of misplaced numbered tiles = 6
- $h_2(N)$ = sum of the (Manhattan) distance of every numbered tile to its goal position
= $2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$

8-Puzzle

$f(N) = h(N) = \text{number of misplaced numbered tiles}$




The white tile is the empty tile

8-Puzzle

$$f(N) = g(N) + h(N)$$

with $h(N)$ = number of misplaced numbered tiles



Heuristic Function

Heuristic

- Use our domain knowledge about the problem to choose some (not all) successors of the current state
- To speed up the searching process
- Heuristic function takes up a state and return the assessment of that state
- Finding right function is not always easy

Heuristic Function and AI search

- Heuristic function estimates how close a state is to the goal state
- It is just an estimation of path cost from a state to goal state (not actual value)
- In state space search, heuristic function helps to reduce the number of nodes expanded during search

Problem relaxation

- Standard approach to create a heuristic is problem relaxation
- Add some new actions for the problem so that search is not required to find the solution cost in the relaxed problem

Problem relaxation with example

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

Goal state

- A problem with fewer restriction on action is called relaxed problem
- One way: pick up any misplaced tiles and place it in the appropriate position (tiles can move anywhere)
- Another way: relaxation on moving tiles along the x-axis and y-axis (tiles can move adjacent squares)

Problem relaxation with example

5		8
4	2	1
7	3	6

STATE(N)

1	2	3
4	5	6
7	8	

Goal state

- $h_1(N)$ = number of misplaced numbered tiles = 6
- $h_2(N)$ = sum of the (Manhattan) distance of every numbered tile to its goal position
= $2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$

Cost of optimal solution to the relaxed problem is an admissible heuristic for the original problem (h_1 and h_2 both are admissible heuristic)

Admissible Heuristic

- It never overestimates the cost for reaching goal node from any node
- Let $h^*(N)$ be the cost of the optimal path from N to a goal node
- The heuristic function $h(N)$ is **admissible** if for all N :

$$0 \leq h(N) \leq h^*(N)$$

$h(N) < h^*(N)$ [underestimated]; $h(N) > h^*(N)$ [overestimated]

- An admissible heuristic function is always **optimistic** !

G is a goal node $\Rightarrow h(G) = 0$

Admissible vs. Non-admissible heuristic


- **Non-admissible heuristic:**

- Pessimistic heuristic because they overestimate the cost
- Breaking optimality by trapping good plan in fringe

- **Admissible heuristic:**

- Optimistic heuristic they can only underestimate the cost
- Can only slow down the search by assigning lower cost to a bad plan so that it will be explored first; but it will find the optimal path gradually.

Inadmissible



$$f(N) = g(N) + h(N)$$

$$1. OL = \{ S^8 \}$$

$$CL = \{ \quad \}$$

$$2. OL = \{ A^7, B^9 \}$$

$$CL = \{ S^8 \}$$

$$3. OL = \{ B^9, G^6 \}$$


$$CL = \{ S^8, A^7 \}$$

$$4. OL = \{ B^9 \}$$

$$CL = \{ S^8, A^7, G^6 \}$$

Break optimality

Admissible



$$1. OL = \{ S^3 \}$$

$$CL = \{ \}$$

$$2. OL = \{ A^4, B^5 \}$$

$$CL = \{ S^3 \}$$

$$3. OL = \{ B^5, G^6 \}$$

$$CL = \{ S^3, A^4 \}$$

$$4. OL = \{ G^5 \}$$

$$CL = \{ S^3, A^4, B^5 \}$$

unnecessary expansion
of A; but find optimal
solution gradually

Dominance

For two admissible heuristics h_1 and h_2

If $h_2(n) \geq h_1(n)$ for all n , then h_2 dominates h_1

h_2 is more informed than h_1 .

h_2 is better for search.

[Better heuristic means we have to explore fewer nodes before finding the solution]

Composite Heuristic

- maximum of two admissible heuristics is also admissible
- Suppose, we have designed two or more heuristics and unsure about that any of them dominates all others
- We can use maximum of them as composite heuristic

$$h(n) = \max \{ h_1(n), \dots, h_m(n) \}$$


Consistent Heuristic

A heuristic h is **consistent** (or **monotone**) if
1) for each node N and each successor N' of N :

$$h(N) \leq w(N, N') + h(N')$$

2) for each goal node G :

$$h(G) = 0$$



A consistent heuristic is also admissible;
but opposite may or may not hold

Consistent Heuristic


- It will be stricter than the admissible one.
- one consequence, f-value never decreases along the path

$$\begin{aligned}f(N') &= g(N') + h(N') \\&= g(N) + w(N, N') + h(N') \\&\geq g(N) + h(N) \\&= f(N) \\f(N') &\geq f(N)\end{aligned}$$


→ Intuition: a consistent heuristic becomes more precise as we get deeper in the search tree

Consistency Violation

If h tells that N is 100 units from the goal, then moving from N along an arc costing 10 units should **not** lead to a node N' that h estimates to be 10 units away from the goal



Consistent and admissible




$$h(N) \leq W(N, N') + h(N')$$

$$f(S) = 4 + 0 = 4$$

$$f(B) = 4 + 1 = 5$$

$$f(G) = 0 + 6 = 6$$

↙
Come closer to
goal state;
evaluation
function increase




Admissible /
no +
Consistent

Best First Search

- It is a way to combining the advantage of both breadth first search and depth first search
- DFS is good if it allows a solution to be found without all competing branches having to be expanded
- BFS is good as it does not trap into a dead end
- Combining by following a single path at a time but switch the path whenever a competing path looks more promising than the current one


Example of Best First Search



Greedy best-first search

- Evaluation function $f(n)=h(n)$, estimated cost from n to goal
- In route finding problem, $h_{SLD}(n)$ =straight line distance/air distance from a node n to destination city
- Heuristic can't be computed from the problem description itself
- In addition, the problem specific domain knowledge is required to understand the correlation between the actual road distance and straight line distance
- The strategy expands the node that appears to be closest to goal (completely heuristic dependant)
- It may lead to the dead end in case of route finding problem

Route finding problem from I to F



start node

goal node

Straight line distance between two cities (h)

Apply Greedy BFS

1. $OL = \{I^{226}\}$

$CL = \{\}$

2. $OL = \{N^{90}, V^{190}\}$

$CL = \{I^{226}\}$

3. $OL = \{V^{190}, I^{226}\}$

$CL = \{I^{226}, N^{90}\}$

Infinite loop (if allow revisit)
otherwise dead end

$h \rightarrow$ estimated one

Evaluation

- If the state space is infinite, in general the search is not complete
- If the state space is finite and we do not discard nodes that revisit states, in general the search is not complete
- If the state space is finite and we discard nodes that revisit states, the search is complete, but in general is not optimal
- The worst case time and space complexity is $O(b^m)$, m is maximum depth of search space
- However, the complexity can be reduced substantially with good heuristic.

A* Search

(most popular algorithm in AI)

1) $f(N) = g(N) + h(N)$, where:

- $g(N)$ = cost of best path found so far to N
- $h(N)$ = heuristic function

2) for all arcs: $w(N, N') \geq \epsilon > 0$

→ Best-first search is then called **A* search**

Search algorithm (A^*)

1. **Initialize:** Set $OPEN = \{s\}$,
 $CLOSED = \{ \}$ Set $g(s) = 0$ and $f(s)=h(s)$
2. **Fail:** If $OPEN = \{ \}$, Terminate & fail
3. **Select:**
Select the minimum cost state, n ,
from $OPEN$ and save n in $CLOSED$
4. **Terminate:**
If $n \in G$, terminate with success

Search algorithm (A^*)

5. Expand:

Generate the successors of n using successor function.

For each successor, m :

If $m \notin [\text{OPEN} \cup \text{CLOSED}]$

Set $g(m) = g(n) + w(n, m)$

Set $f(m) = g(m) + h(m)$

and insert m in OPEN

If $m \in [\text{OPEN} \cup \text{CLOSED}]$

Set $g(m) = \min \{g(m), g(n) + w(n, m)\}$


Set $f(m) = g(m) + h(m)$

If $f(m)$ has decreased and

$m \in \text{CLOSED}$, move it to OPEN

Evaluation


- A* search is not optimal if the heuristic is overestimated (break optimality)
- A* search is optimal with admissible heuristic if node reopening is allowed.
- Node reopening leads to the unnecessary node expansion (re-visit)
- A* search is optimal with consistent heuristic; In this case, node reopening is not needed (always expand node with optimal path)



Admissible heuristic

$$h(n) \leq h^*(n)$$

1. $OL = \{S^{15}\}$ $CL = \{\}$
2. $OL = \{A^{14}, B^{28}\}$ $CL = \{S^{15}\}$
 due to accurate heuristic, good plan trapped
3. $OL = \{B^{28}, D^{23}\}$ $CL = \{S^{15}, A^{14}\}$
 trap and node in bad plan
4. $OL = \{B^{28}, E^{34}\}$ $CL = \{S^{15}, A^{14}, D^{23}\}$
 move from closed to open
5. $OL = \{E^{34}, D^{19}\}$ $CL = \{S^{15}, A^{14}, D^{23}, B^{28}\}$
 how we have a good plan
6. $OL = \{E^{34}, 30\}$ $CL = \{S^{15}, A^{14}, D^{19}, E^{30}\}$
 discard previous path
7. $OL = \{G^{30}\}$ $CL = \{S^{15}, A^{14}, B^{28}, D^{19}, E^{30}\}$
 we can reach optimal good in next step



Consistent heuristic

$$h(n) \leq w(n, n') + h(n')$$

1. OL = { S¹³ }, CL = { }

2. OL = { A¹⁴, B¹⁷ }
CL = { S¹³ }

3. OL = { B¹⁷, D²³ }
CL = { S¹³, A¹⁴ }

4. OL = { D²³ }
CL = { S¹³, A¹⁴, B¹⁷ }

CL = { S¹³, A¹⁴, B¹⁷ } → good plan

5. OL = { E³⁰ }

CL = { S¹³, A¹⁴, B¹⁷, D¹⁹ }

6. OL = { G³⁰ }

CL = { S¹³, A¹⁴, B¹⁷, D¹⁹, E³⁰ } → goal in next step

Completeness

- A^* expands all nodes with $f(n) < C^*$ (C^* is the cost of optimal solution path)
- It may expand some nodes with $f(n) = C^*$
- A^* search is complete if there is finitely many nodes with cost less than or equal to C^* (additionally step cost exceeds some finite value and b is finite)

Complexity

- Time/space complexity is exponential
- With good heuristic, it can reduce significantly
- Main problem is storage; need to store all generated nodes

A* Search (optimally efficient)

- A* search is optimally efficient for any given consistent heuristic
- No other optimal algorithm is guaranteed to expand fewer nodes than A*
- Any algorithm that does not expand all node with $f(n) < C^*$ runs the risk of missing the optimal solution.

Conclusion

- Design good heuristic for A^* search
- Use different variants of A^* (if possible some bound on memory requirement)


We can think about a solution (not necessarily optimal) for large-scale AI problem.

Move to the local search algorithm


Local Search

- Light-memory search method
- No search tree; only the current state is represented!
- Only applicable to problems where the path is irrelevant (e.g., 8-queen)
- Many similarities with optimization techniques


Search problems



Blind search




Heuristic search:
best-first (A^* and Greedy BF)



Construction of heuristics



Variants of A^*



Local search