

Chapter 3: Introduction to SQL

Edited by Radhika Sukapuram. Original slides by


Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Schema Diagram for University Database





Natural Join cont.

- List the names of instructors along with the courses that they teach

```
select name, title  
from instructor natural join teaches, course  
where teaches.course_id = course.course_id;
```

Does the query below compute the same result ?

```
select name, title  
from instructor natural join teaches natural join course
```





Natural Join cont.

- List the names of instructors along with the courses that they teach

```
select name, title  
from instructor natural join teaches, course  
where teaches.course_id = course.course_id;
```

Does the query below compute the same result ?

```
select name, title  
from instructor natural join teaches natural join course
```

No!

To avoid equating attributes wrongly,

```
select name, title  
from (instructor natural join teaches) join course using  
(course_id);
```



The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

old-name as new-name

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

- **select distinct** *T.name*
from *instructor as T, instructor as S*
where *T.salary > S.salary and S.dept_name = 'Comp. Sci.'*

- Keyword **as** is optional and may be omitted
instructor as T \equiv instructor T

Note: T is called a correlation name, table alias, correlation variable or tuple variable.



String Operations

- ❑ SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
 - ❑ percent (%). The % character matches any substring.
 - ❑ underscore (_). The _ character matches any character.
- ❑ Find the names of all instructors whose name includes the substring “dar”.

```
select name  
from instructor  
where name like '%dar%'
```

- ❑ Match the string “100%”

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.



String Operations (Cont.)

- ❑ Patterns are case sensitive.
- ❑ Pattern matching examples:
 - ❑ 'Intro%' matches any string beginning with "Intro".
 - ❑ '%Comp%' matches any string containing "Comp" as a substring.
 - ❑ '___' matches any string of exactly three characters.
 - ❑ '___ %' matches any string of at least three characters.
- ❑ SQL supports a variety of string operations such as
 - ❑ concatenation (using "||")
 - ❑ converting from upper to lower case (and vice versa)
 - ❑ finding string length, extracting substrings, etc.



Ordering the Display of Tuples

- List in alphabetic order the names of all instructors

```
select distinct name  
from instructor  
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
 - Example: **order by** *name* **desc**
- Can sort on multiple attributes
 - Example: **order by** *dept_name*, *name*



Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, $\geq \$90,000$ and $\leq \$100,000$)
 - **select** *name*
from *instructor*
where *salary* **between** 90000 **and** 100000
- Tuple comparison
 - **select** *name, course_id*
from *instructor, teaches*
where (*instructor.ID, dept_name*) = (*teaches.ID, 'Biology'*);



section

<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>	<i>building</i>	<i>room_number</i>	<i>time_slot_id</i>
BIO-101	1	Summer	2009	Painter	514	B
BIO-301	1	Summer	2010	Painter	514	A
CS-101	1	Fall	2009	Packard	101	H
CS-101	1	Spring	2010	Packard	101	F
CS-190	1	Spring	2009	Taylor	3128	E
CS-190	2	Spring	2009	Taylor	3128	A
CS-315	1	Spring	2010	Watson	120	D
CS-319	1	Spring	2010	Watson	100	B
CS-319	2	Spring	2010	Taylor	3128	C
CS-347	1	Fall	2009	Taylor	3128	A
EE-181	1	Spring	2009	Taylor	3128	C
FIN-201	1	Spring	2010	Packard	101	B
HIS-351	1	Spring	2010	Painter	514	C
MU-199	1	Spring	2010	Packard	101	D
PHY-101	1	Fall	2009	Watson	100	A



Set Operations

- Find courses that ran in Fall 2009 or in Spring 2010

- Find courses that ran in Fall 2009 and in Spring 2010

- Find courses that ran in Fall 2009 but not in Spring 2010



Set Operations

- Find courses that ran in Fall 2009 or in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)  
union  
(select course_id from section where sem = 'Spring' and year = 2010)
```

- Find courses that ran in Fall 2009 and in Spring 2010

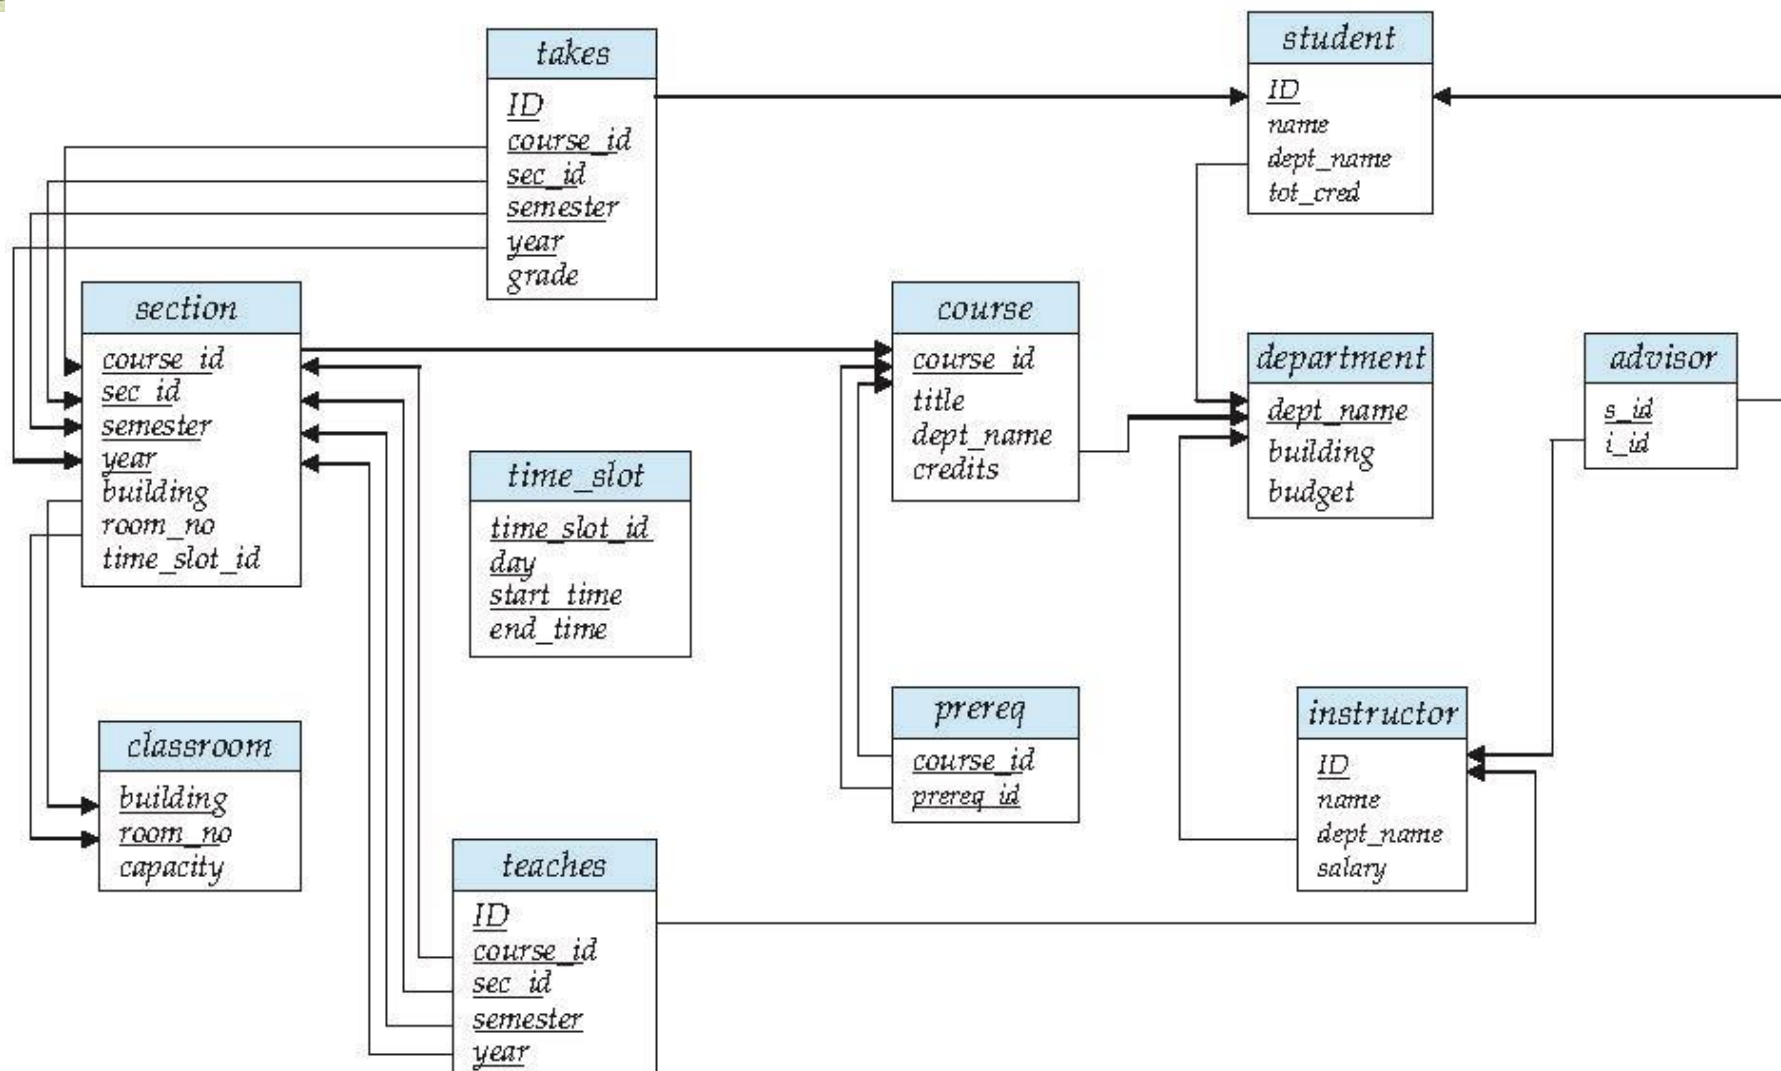
```
(select course_id from section where sem = 'Fall' and year = 2009)  
intersect  
(select course_id from section where sem = 'Spring' and year = 2010)
```

- Find courses that ran in Fall 2009 but not in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)  
except  
(select course_id from section where sem = 'Spring' and year = 2010)
```



Set Operations





Set Operations (Cont.)

- Find the largest salary of all instructors.
 - ▶ Find the salaries of all instructors that are less than the largest salary.
 - ▶ Find all the salaries of all instructors
 - ▶ Find the largest salary of all instructors.





Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary.
 - **select distinct** *T.salary*
from *instructor* **as** *T*, *instructor* **as** *S*
where *T.salary* < *S.salary*

- Find all the salaries of all instructors
 - **select distinct** *salary*
from *instructor*

- Find the largest salary of all instructors.
 - **(select** “second query”)
except
(select “first query”)



Set Operations (Cont.)

- Set operations **union**, **intersect**, and **except**
 - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs m times in r and n times in s , then, it occurs:
 - $m + n$ times in r **union all** s
 - $\min(m, n)$ times in r **intersect all** s
 - $\max(0, m - n)$ times in r **except all** s



Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of the attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
 - Example: $5 + \text{null}$ returns null
- The predicate **is null** can be used to check for null values.
 - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```



Null Values and Three Valued Logic

- Three values – *true*, *false*, *unknown*
- Any comparison with *null* returns *unknown*
 - Example: $5 < null$ or $null <> null$ or $null = null$
- Three-valued logic using the value *unknown*:
 - OR: $(unknown \text{ or } true) = true$,
 $(unknown \text{ or } false) = unknown$
 $(unknown \text{ or } unknown) = unknown$
 - AND: $(true \text{ and } unknown) = unknown$,
 $(false \text{ and } unknown) = false$,
 $(unknown \text{ and } unknown) = unknown$
 - NOT: $(\text{not } unknown) = unknown$
 - “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



Where predicate with null

- Example: Find all instructors whose salary is greater than 5000.

```
select name  
from instructor  
where salary > 500
```

Suppose a record has salary = null.

What value will the predicate return for that record ?

Will it be treated as true or false ?



Where predicate with null

- Example: Find all instructors whose salary is greater than 5000.

```
select name  
from instructor  
where salary > 500
```

Suppose a record has salary = null.

What value will the predicate return for that record ?

Will it be treated as true or false ?

As false



null: select distinct, set operations

- Eliminating duplicate tuples

```
select distinct name  
from instructor  
where salary > 500
```

Suppose two records have name as null

Will both the records be retained ?



null: select distinct, set operations

- Eliminating duplicate tuples

```
select distinct name  
from instructor  
where salary > 500
```

Suppose two records have name as null

Will both the records be retained ?

No, only one will be retained – treatment different from predicates

- Same treatment for set operations