

Chapter 3: Introduction to SQL

Edited by Radhika Sukapuram. Original slides by

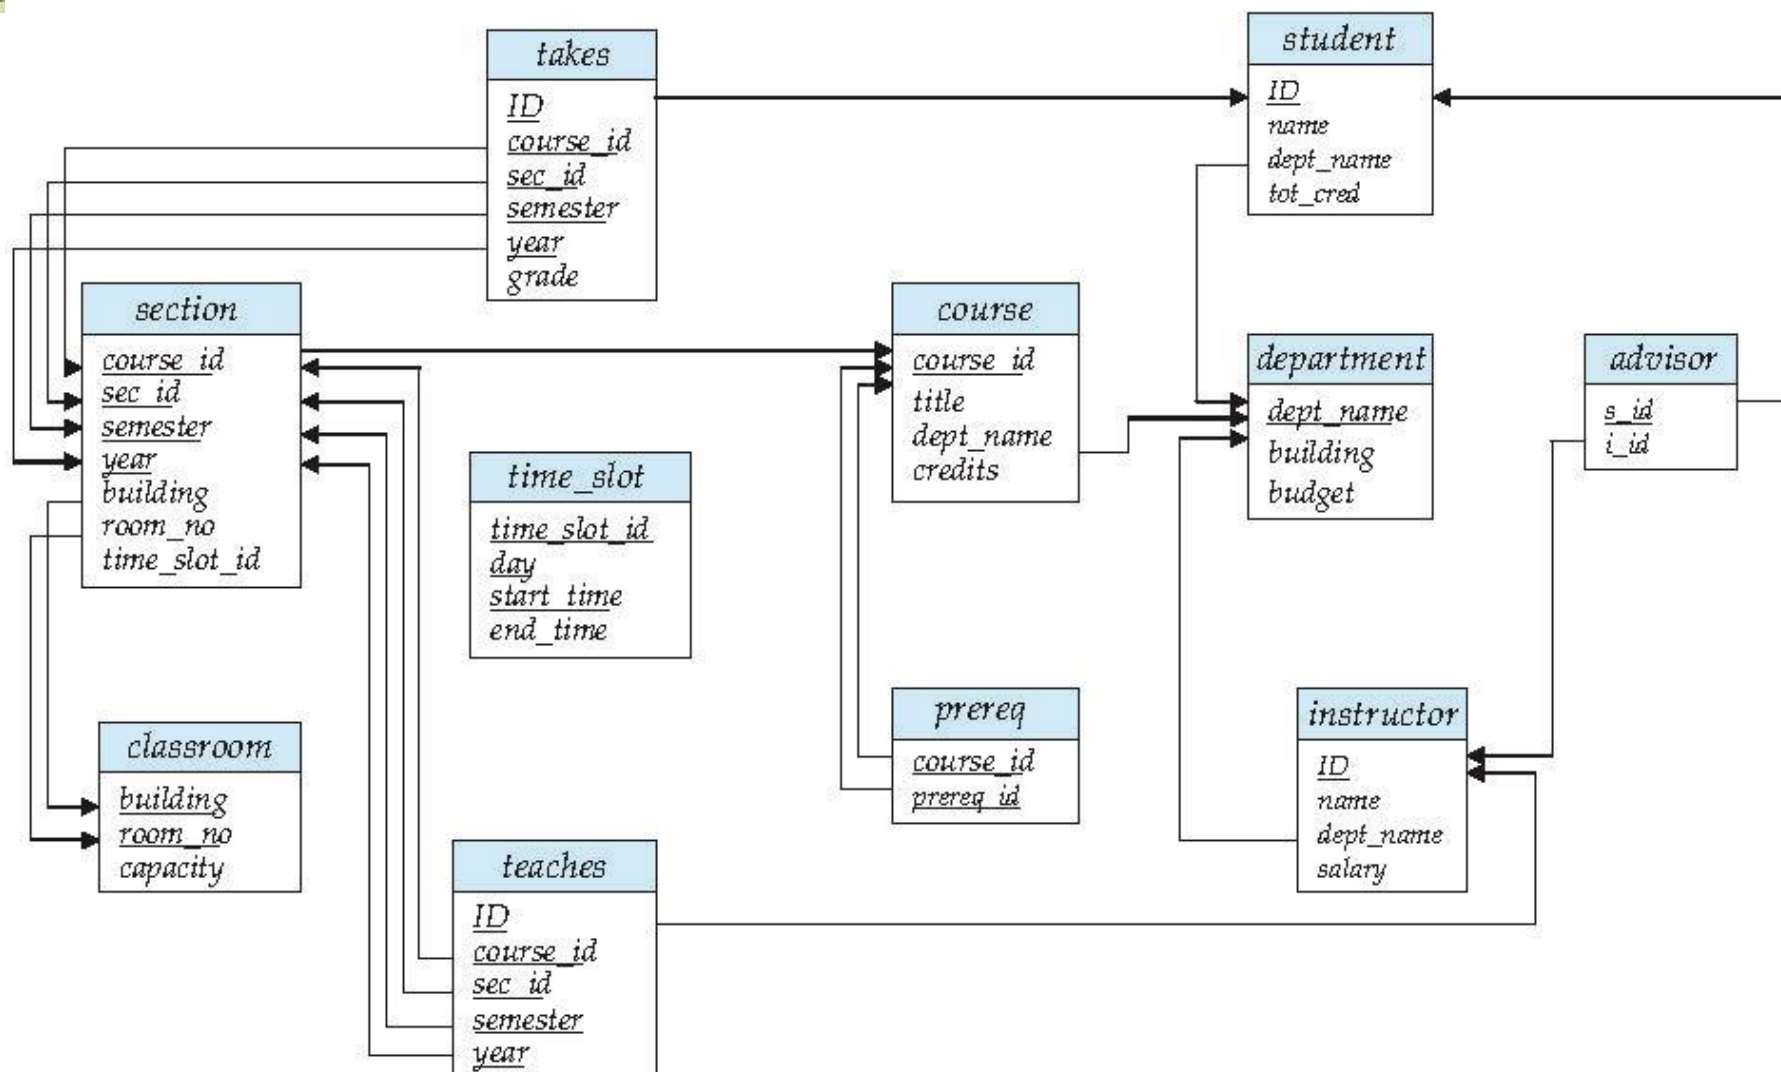
Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Set Operations





Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary.
- Find all the salaries of all instructors
- Find the largest salary of all instructors.



Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary.
 - **select distinct** *T.salary*
from *instructor* **as** *T*, *instructor* **as** *S*
where *T.salary* < *S.salary*

- Find all the salaries of all instructors
 - **select distinct** *salary*
from *instructor*

- Find the largest salary of all instructors.
 - (**select** “second query”)
except
(**select** “first query”)



Example

Instructor
<u>Salary</u>
5
8
1



T
<u>Salary</u>
5
8
1

S
<u>Salary</u>
5
8
1



Temp	
T	S
<u>Salary</u>	<u>Salary</u>
5	8
1	5
1	8

T	S
<u>Salary</u>	<u>Salary</u>
5	5
5	8
5	1
8	5
8	8
8	1
1	5
1	8
1	1





Temp	
T	S
<u>Salary</u>	<u>Salary</u>
5	8
1	5
1	8



Instructor
<u>Salary</u>
5
8
1

−

Temp
<u>T.Salary</u>
5
1

=

Result
<u>Salary</u>
8



Set Operations (Cont.)

- Set operations **union**, **intersect**, and **except**
 - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs m times in r and n times in s , then, it occurs:
 - $m + n$ times in r **union all** s
 - $\min(m, n)$ times in r **intersect all** s
 - $\max(0, m - n)$ times in r **except all** s



Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of the attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
 - Example: $5 + \text{null}$ returns null
- The predicate **is null** can be used to check for null values.
 - Example: Find all instructors whose salary is null.

```
select name
from instructor
where salary is null
```




Null Values and Three Valued Logic

- Three values – *true*, *false*, *unknown*
- Any comparison with *null* returns *unknown*
 - Example: $5 < null$ or $null <> null$ or $null = null$
- Three-valued logic using the value *unknown*:
 - OR: $(unknown \text{ or } true) = true$,
 $(unknown \text{ or } false) = unknown$
 $(unknown \text{ or } unknown) = unknown$
 - AND: $(true \text{ and } unknown) = unknown$,
 $(false \text{ and } unknown) = false$,
 $(unknown \text{ and } unknown) = unknown$
 - NOT: $(\text{not } unknown) = unknown$
 - “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



Where predicate with null

- Example: Find all instructors whose salary is greater than 5000.

```
select name  
from instructor  
where salary > 5000
```

Suppose a record has salary = null.

What value will the predicate return for that record ?

Will it be treated as true or false ?



Where predicate with null

- Example: Find all instructors whose salary is greater than 5000.

```
select name  
from instructor  
where salary > 500
```

Suppose a record has salary = null.

What value will the predicate return for that record ?

Will it be treated as true or false ?

As false



null: select distinct, set operations

- Eliminating duplicate tuples

```
select distinct name  
from instructor  
where salary > 500
```

Suppose two records have name as null

Will both the records be retained ?



null: select distinct, set operations

- Eliminating duplicate tuples

```
select distinct name  
from instructor  
where salary > 500
```

Suppose two records have name as null

Will both the records be retained ?

No, only one will be retained – treatment different from predicates

- Same treatment for set operations



Aggregate Functions

- These functions operate on the (collection) multiset of values of a column of a relation, and return a value

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values



instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

teaches

<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009



Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department
 - **select avg** (*salary*)
from *instructor*
where *dept_name*= 'Comp. Sci.';
- Find the total number of instructors who teach a course in the Spring 2010 semester
 - **select count** (**distinct** *ID*)
from *teaches*
where *semester* = 'Spring' **and** *year* = 2010;
- Find the number of tuples in the *course* relation
 - **select count** (*)
from *course*;



instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000



Aggregate Functions – Group By

- Find the average salary of instructors in each department

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000



Aggregate Functions – Group By

- Find the average salary of instructors in each department
 - **select** *dept_name*, **avg** (*salary*) **as** *avg_salary*
from *instructor*
group by *dept_name*;

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

<i>dept_name</i>	<i>avg_salary</i>
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000



Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
 - */* erroneous query */*
select *dept_name, ID, avg (salary)*
from *instructor*
group by *dept_name;*
 - Which id must be output corresponding to a group (dept_name) is not known



Aggregate Functions – Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups



Order of evaluation of aggregate queries

select A_1, A_2, \dots, A_n	5
from r_1, r_2, \dots, r_m	1
where P	2
group by A_1, A_2, \dots, A_n	3
having P_2	4

1. **from** clause is evaluated first to get a relation
2. If a **where** clause is present, its predicate is applied
3. Tuples are then grouped as per **group by** if it exists
4. **having** clause is applied to each group; groups that do not satisfy the clause are removed
5. **select** clause uses tuples of the remaining groups to generate the result



Null Values and Aggregates

- Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if all amounts are null
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
 - count returns 0
 - all other aggregates return null